



UNIVERSITÉ
CÔTE D'AZUR

Lecture 5: First contact with Transformers

Advanced deep learning

Rémy Sun

remy.sun@inria.fr

inria



Course organization


- Goal: In-depth understanding of important Deep Learning staples
 - Reinforce what you have already seen
 - Introduce state of the art models
- This is a hands-on course in pytorch
 - Minimal math
 - Enough to understand
 - Quite a bit of coding
 - Get comfortable with the standard pipeline

- Hand in one or two lab notebooks
 - Questions + (clean) code
 - 1st notebook: Lab 5 on transformers (Today!)
 - To hand in after vacation (Nov. 4 AoE)
- Written exam at end of semester
 - Little to no code
 - A few exercises on toy examples
 - Questions on aspects of deep learning

- L1-2: Overview of Deep Learning (F. Precioso)
- L3-4: Fundamentals of Deep Learning (R. Sun)
- **L5-6: Transformers (R. Sun)**
- L7: Large models (LLMs, VLMs, Generators) (R. Sun)
- L8: Tricks of the trade (R. Sun)
- L9: Ethics of AI (F. Precioso)
- L10: Intro to generative models (P-A. Mattei)

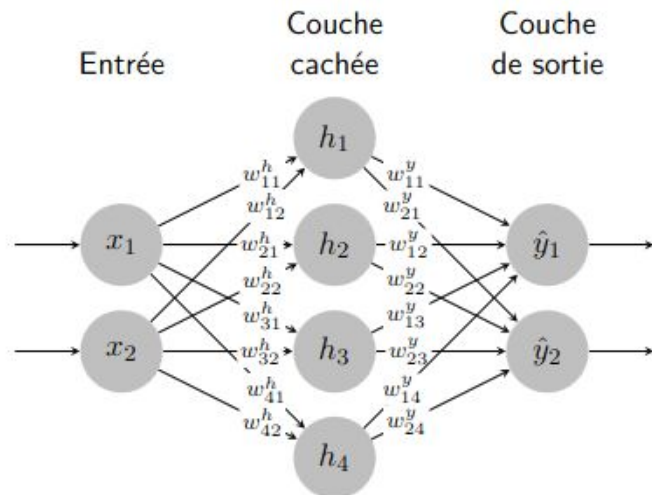
- Goal: Understand basic deep architectures in-depth
 - Building blocks for everything else in deep learning
 - Deep learning relies on the combination of a lot of very simple blocks
- A few things to take away after these 3 lectures
 - What do we optimize for? How? Why?
 - How do we build and train neural layers?
 - How do they behave?

Refresher on last week



- Simple 1 hidden layer MLP
 - 2 inputs
 - 2 outputs
 - 4 hidden activations
- Classification problem
 - Outputs probabilities
 - Cross-entropy loss

$$l_{CE}(\hat{y}, y) = - \sum_{i=0}^{\#Classes-1} y_i \log(\hat{y}_i)$$



$$\begin{cases} \tilde{\mathbf{h}} = \mathbf{x}\mathbf{W}^h + \mathbf{b}^h \\ \mathbf{h} = \tanh(\tilde{\mathbf{h}}) \\ \tilde{\mathbf{y}} = \mathbf{h}\mathbf{W}^y + \mathbf{b}^y \\ \hat{\mathbf{y}} = \text{SoftMax}(\tilde{\mathbf{y}}) \end{cases}$$

$$\begin{cases} \nabla_{\hat{\mathbf{y}}} = \hat{\mathbf{y}} - \mathbf{y} \\ \nabla_{\mathbf{W}^y} = \nabla_{\hat{\mathbf{y}}}^\top \mathbf{h} \\ \nabla_{\mathbf{b}^y} = \nabla_{\hat{\mathbf{y}}}^\top \\ \nabla_{\tilde{\mathbf{h}}} = (\nabla_{\hat{\mathbf{y}}} \mathbf{W}^y) \odot (1 - \mathbf{h}^2) \\ \nabla_{\mathbf{W}^h} = \nabla_{\tilde{\mathbf{h}}}^\top \mathbf{x} \\ \nabla_{\mathbf{b}^h} = \nabla_{\tilde{\mathbf{h}}}^\top \end{cases}$$

- Lecture 1 practical correction on Moodle
 - Implement this by hand with basic torch!
 - Careful with batch dimension!

```
1 def init_params(nx, nh, ny):
2     """
3     nx, nh, ny: integers
4     out params: dictionary
5     """
6     params = {}
7
8     params["Wh"] = torch.randn((nh, nx))*0.3
9     params["Wy"] = torch.randn((ny, nh))*0.3
10    params["bh"] = torch.zeros((nh,1))
11    params["by"] = torch.zeros((ny,1))
12
13    return params
```

```
1 def forward(params, X):
2     """
3     params: dictionnary
4     X: (n_batch, dimension)
5     """
6     bsize = X.size(0)
7     nh = params['Wh'].size(0)
8     ny = params['Wy'].size(0)
9     outputs = {}
10
11     outputs["X"] = X
12     outputs["htilde"] = torch.mm(X, params["Wh"].T) + params["bh"].T
13     outputs["h"] = torch.tanh(outputs["htilde"])
14     outputs["ytilde"] = torch.mm(outputs["h"], params["Wy"].T) + params["by"].T
15     outputs["yhat"] = torch.exp(outputs["ytilde"])
16     outputs["yhat"] = outputs["yhat"] / outputs["yhat"].sum(dim=-1, keepdim=True)
17
18
19     return outputs['yhat'], outputs
```

```
1 def backward(params, outputs, Y):
2     bsize = Y.shape[0]
3     grads = {}
4
5     Y_tilde_grad = outputs["yhat"] - Y
6     h_tilde_grad = torch.mm(Y_tilde_grad, params['Wy']
7                             ) * (1 - torch.pow(outputs['h'], 2))
8
9     grads["Wy"] = torch.mm(Y_tilde_grad.T, outputs["h"])
10    grads["Wh"] = torch.mm(h_tilde_grad.T, outputs['X'])
11    grads["by"] = Y_tilde_grad.sum(dim=0, keepdim=True).T
12    grads["bh"] = h_tilde_grad.sum(0, keepdim=True).T
13
14    grads['Wy'] /= bsize
15    grads['by'] /= bsize
16    grads['Wh'] /= bsize
17    grads['bh'] /= bsize
18
19    return grads
```

```
1 def sgd(params, grads, eta):  
2  
3     params['Wy'] -= eta * grads['Wy']  
4     params['Wh'] -= eta * grads['Wh']  
5     params['by'] -= eta * grads['by']  
6     params['bh'] -= eta * grads['bh']  
7  
8     return params
```

```
for j in range(N // Nbatch):  
  
    indsBatch = range(j * Nbatch, (j+1) * Nbatch)  
    X = Xtrain[indsBatch, :]  
    Y = Ytrain[indsBatch, :]  
  
    Y_hat, outputs = forward(params, X)  
    loss, accuracy = loss_accuracy(Y_hat, Y)  
    grads = backward(params, outputs, Y)  
    params = sgd(params, grads, eta)
```

```
def backward(params, outputs, Y):
    bsize = Y.shape[0]
    grads = {}

    Y_tilde_grad = outputs["yhat"] - Y
    h_tilde_grad = torch.mm(Y_tilde_grad, params['Wy'])
    h_tilde_grad = h_tilde_grad * (1 - torch.pow(outputs['h'], 2))

    grads["Wy"] = torch.mm(Y_tilde_grad.T, outputs["h"])
    grads["Wh"] = torch.mm(h_tilde_grad.T, outputs["X"])
    grads["by"] = Y_tilde_grad.sum(dim=0, keepdim=True).T
    grads["bh"] = h_tilde_grad.sum(0, keepdim=True).T

    grads['Wy'] /= bsize
    grads['by'] /= bsize
    grads['Wh'] /= bsize
    grads['bh'] /= bsize

    return grads
```

- Torch.tensor object
 - Np.array like
 - Tracked on a computational graph
 - .grad variable to track gradients
 - .backward to backpropagate gradients through the graph
 - Activate .autograd!


```
def backward(params, outputs, Y):
    bsize = Y.shape[0]
    grads = {}

    Y_tilde_grad = outputs["yhat"] - Y
    h_tilde_grad = torch.mm(Y_tilde_grad, params['Wy'])
                        * (1 - torch.pow(outputs['h'], 2))

    grads["Wy"] = torch.mm(Y_tilde_grad.T, outputs["h"])
    grads["Wh"] = torch.mm(h_tilde_grad.T, outputs["X"])
    grads["by"] = Y_tilde_grad.sum(dim=0, keepdim=True).T
    grads["bh"] = h_tilde_grad.sum(0, keepdim=True).T

    grads['Wy'] /= bsize
    grads['by'] /= bsize
    grads['Wh'] /= bsize
    grads['bh'] /= bsize

    return grads
```

```
params['Wh'] = torch.randn(nh, nx) * 0.3
params['Wh'].requires_grad = True
params['bh'] = torch.zeros(nh, 1, requires_grad=True)
params['Wy'] = torch.randn(ny, nh) * 0.3
params['Wy'].requires_grad = True
params['by'] = torch.zeros(ny, 1, requires_grad=True)
```

```
with torch.no_grad():
    params['Wy'] -= eta * params['Wy'].grad
    params['Wh'] -= eta * params['Wh'].grad
    params['by'] -= eta * params['by'].grad
    params['bh'] -= eta * params['bh'].grad

    params['Wy'].grad.zero_()
    params['Wh'].grad.zero_()
    params['by'].grad.zero_()
    params['bh'].grad.zero_()
```

```
Yhat, outputs = forward(params, X)
L, _ = loss_accuracy(Yhat, Y)
L.backward()
params = sgd(params, 0.03)
```


- Torch keeps track of operations in a dynamic graph
 - Not static like tensorflow
 - Need to tell pytorch to do it
 - `.requires_grad=True`
 - Have to exclude graph operations
 - With `torch.no_grad()`:
- Gradients are automatically computed
 - `L.backward()`
 - `L` is a scalar loss tensor
 - Go back through the graph to fill `.grad`

```
params = {}

params["Wh"] = torch.randn((nh, nx))*0.3
params["Wy"] = torch.randn((ny, nh))*0.3
params["bh"] = torch.zeros((nh,1))
params["by"] = torch.zeros((ny,1))
```

```
def forward(params, X):
    """
    params: dictionary
    X: (n_batch, dimension)
    """
    bsize = X.size(0)
    nh = params['Wh'].size(0)
    ny = params['Wy'].size(0)
    outputs = {}

    outputs["X"] = X
    outputs["htilde"] = torch.mm(X, params["Wh"].T) + params["bh"].T
    outputs["h"] = torch.tanh(outputs["htilde"])
    outputs["ytilde"] = torch.mm(outputs["h"], params["Wy"].T) + params["by"].T
    outputs["yhat"] = torch.exp(outputs["ytilde"])
    outputs["yhat"] = outputs["yhat"] / outputs["yhat"].sum(dim=-1, keepdim=True)

    return outputs['yhat'], outputs
```

- Torch.nn.Module objects
 - `__init__` creates weights and initializes them!
 - `.forward` implements forward operations
 - Default object call
 - `model(x)`
 - Some global control over model weights

```
params = {}

params["Wh"] = torch.randn((nh, nx))*0.3
params["Wy"] = torch.randn((ny, nh))*0.3
params["bh"] = torch.zeros((nh,1))
params["by"] = torch.zeros((ny,1))
```

```
def forward(params, X):
    """
    params: dictionnary
    X: (n_batch, dimension)
    """
    bsize = X.size(0)
    nh = params['Wh'].size(0)
    ny = params['Wy'].size(0)
    outputs = {}

    outputs["X"] = X
    outputs["htilde"] = torch.mm(X, params["Wh"].T) + params["bh"].T
    outputs["h"] = torch.tanh(outputs["htilde"])
    outputs["ytilde"] = torch.mm(outputs["h"], params["Wy"].T) + params["by"].T
    outputs["yhat"] = torch.exp(outputs["ytilde"])
    outputs["yhat"] = outputs["yhat"] / outputs["yhat"].sum(dim=-1, keepdim=True)

    return outputs['yhat'], outputs
```

```
model = torch.nn.Sequential(
    torch.nn.Linear(nx, nh),
    torch.nn.Tanh(),
    torch.nn.Linear(nh, ny)
)
loss = torch.nn.CrossEntropyLoss()
```

```
_, indY = torch.max(Y, 1)
L = loss(Yhat, indY)

_, indYhat = torch.max(Yhat, 1)

acc = torch.sum(indY == indYhat.data) * 100 // indY.size(0);
```

```
with torch.no_grad():
    for param in model.parameters():
        param -= eta * param.grad
    model.zero_grad()
```

```
Yhat = model(X)
L, _ = loss_accuracy(loss, Yhat, Y)
L.backward()
model = sgd(model, 0.03)
```

- What is a network?
 - Weights to be learned
 - A forward function that uses those weights
- Torch.nn modules implement this
 - Initialize weights in `__init__`
 - Implement forward
 - Default call of the object: `model(x)`
 - Large panel of existing layer implementations!
 - Even losses can be made as modules

```
def sgd(params, grads, eta):  
  
    params['Wy'] -= eta * grads['Wy']  
    params['Wh'] -= eta * grads['Wh']  
    params['by'] -= eta * grads['by']  
    params['bh'] -= eta * grads['bh']  
  
    return params
```

- Torch.optim objects
 - Tracks learning rates
 - Tracks weights to optimize
 - Performs SGD steps
 - Even cleans up!

```
def sgd(params, grads, eta):  
  
    params['Wy'] -= eta * grads['Wy']  
    params['Wh'] -= eta * grads['Wh']  
    params['by'] -= eta * grads['by']  
    params['bh'] -= eta * grads['bh']  
  
    return params
```

```
optim = torch.optim.SGD(model.parameters(), lr=eta)
```


```
yhat = model(X)  
L, acc = loss_accuracy(loss, yhat, Y)  
optim.zero_grad()  
L.backward()  
optim.step()
```

- What is an optimizer?
 - A way to update gradients
 - Operation on gradients
 - With parameters
 - (Possibly dynamic)
- Torch.optim optimizer
 - Instantiate as an object
 - With weights to track
 - With parameters of the optimization
 - Call `optimizer.zero_grad()` and `optimizer.step()`


- Training a network requires
 - Weights
 - A forward function
 - A backward function
 - Gradient steps
- Nice pytorch tools
 - `Torch.tensor` and `torch.autograd`
 - `Torch.nn`
 - `Torch.optim`


- Create neural network
 - Use the torch.nn modules
 - Can use torch.nn.Modules
 - Or create a new class inheriting from torch.nn modules
- In training loop

```
yhat = model(X)
L, acc = loss_accuracy(loss, yhat, Y)
optim.zero_grad()
L.backward()
optim.step()
```





- Convolutional layers
 - Local correlations
 - Well suited to images
 - Used sometimes with Transformers now






- Classical architecture of computer vision
 - Convolutional layers for feature extraction
 - Dense/linear layers to make decisions from features
 - E.g. LeNet5 (Before 2000!)


- A few milestones
 - Load image data
 - Load a classic model
 - Train it!
 - (Finetune a strong model)
- Apply knowledge from TP4a!




- Recurrent networks
 - Temporal correlations
 - How to take the past into account?
 - Recent resurgence (SSMs, ...)




one to one




one to many




many to one




many to many




many to many

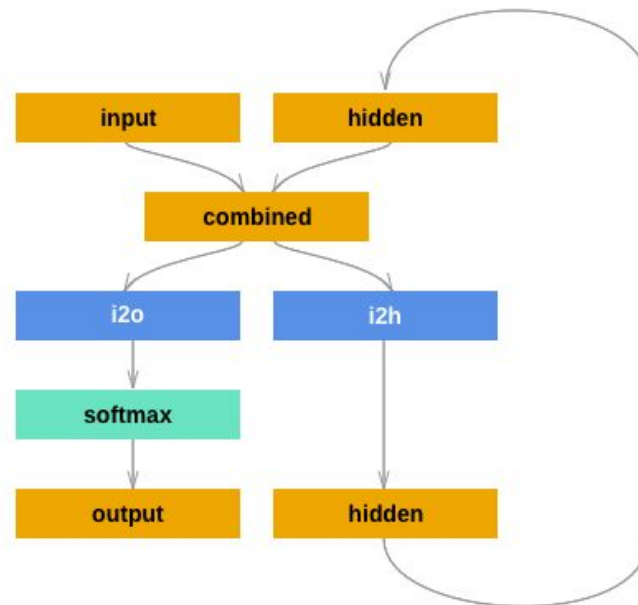




- Basic RNN model
 - Input + hidden state
 - Hidden state remains from input to input



- Making a language processor
 - Tokenize words
 - Create network
 - Train network
 - Apply network
- Apply knowledge from TP4a!



```
class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(RNN, self).__init__()

        self.input_size = input_size
        self.hidden_size = hidden_size

        #####
        ## Your code here ##
        #####

        # Init hidden state

        # Create Linear layers
        self.i2o = None
        self.i2h = None
        # Create activation function
        self.activation = None

    def forward(self, input):

        # Concatenate input and current hidden state
        concat = None

        # Pass through the two hidden layers the concatenation
        output = None
        hidden = None


        # Save the hidden state for later
        self.hidden = None

        # Compute the output
        output = None

        #####
        ## FIN ##
        #####

        return output

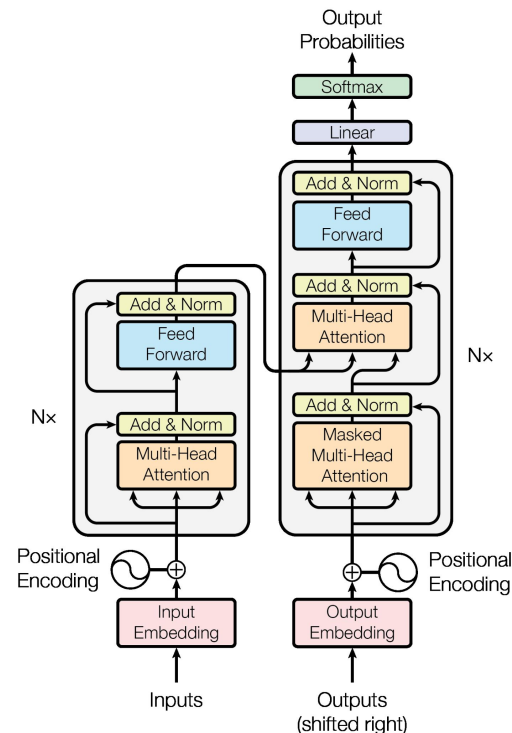
    def init_hidden(self): # to reset hidden state
        self.hidden = torch.zeros(1, self.hidden_size)
```




Transformers: A new neural architecture

- What is the state of the art in
 - Computer vision?
 - CNNs
 - Natural language processing?
 - RNNs
 - Time series?
 - RNNs or TCNs
 - Multimodal problems?
 - Hybrid?

- What is the state of the art in
 - Computer vision?
 - ~~CNNs~~ -> **Transformers**
 - Natural language processing?
 - ~~RNNs~~ -> **Transformers**
 - Time series?
 - ~~RNNs or TCNs~~ -> **Transformers**
 - Multimodal problems?
 - ~~Hybrid?~~ -> **Transformers**
- Transformers use keeps increasing over time







(a) Dense layer

$$d_{\theta}(x) = \sigma(W_{\theta}x^T + b_{\theta})$$

$$\sigma(x) = \text{ReLU}(x) = \max(0, x)$$






- RNNs have issues with long term memory
 - Try to allow it to look at past hidden states
 - That is a lot
 - **Attention** as a solution (~2015)
 - Only look at some past states
- Attention is all you need (Vaswani et al. 2017)
 - **The** transformer
 - Only attention and dense layers
 - Outperforms RNNs substantially

- Attention is all you need (Vaswani et al. 2017)
 - **The** transformer
 - Only attention and dense layers
 - Outperforms RNNs substantially
- It gets harder in Computer Vision
 - Multiple attempts since 2017
 - An Image is Worth 16x16 Words (DeSovitskyi et al. 2021)
 - (First) Vision Transformer that kind of works
 - With all the data in the world
 - Data efficient Image Transformers (Touvron et al. 2021)
 - Works with reasonable datasets

Attention introduction

Dog



Garden



- What is a Dog?
 - It is a 4 legged animal with fur and ears and eyes and a head and ...
 - It is on this part of that picture.


The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .
The FBI is chasing a criminal on the run .

- Define a word by other words it relates to



- Define a word by a few similar pictures

Attention mechanism



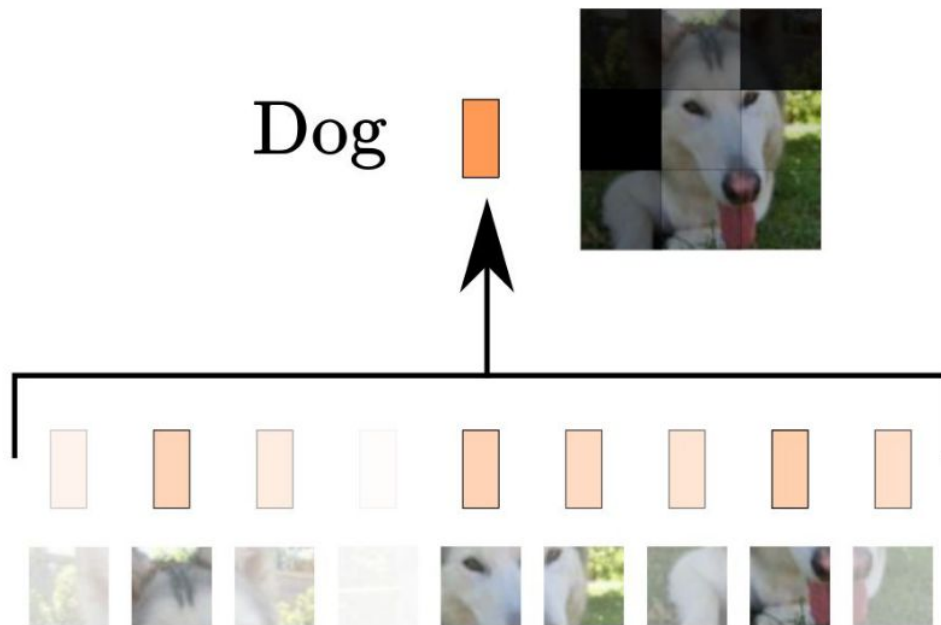
- Generic example between text and image
 - Could be something else!
- How do you define the dog with the picture?



- Find the relevant parts of the picture

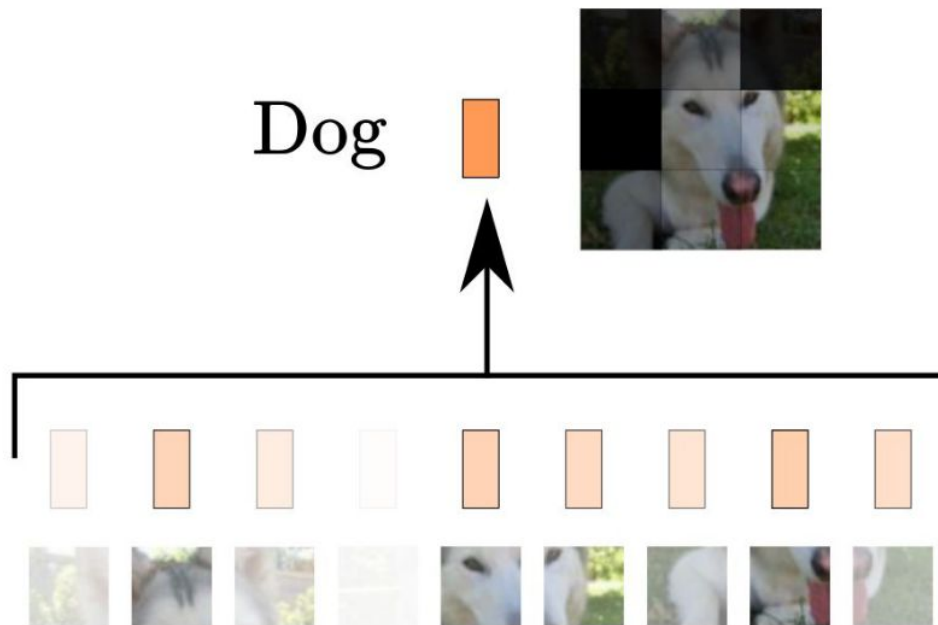


- Find the relevant parts of the picture

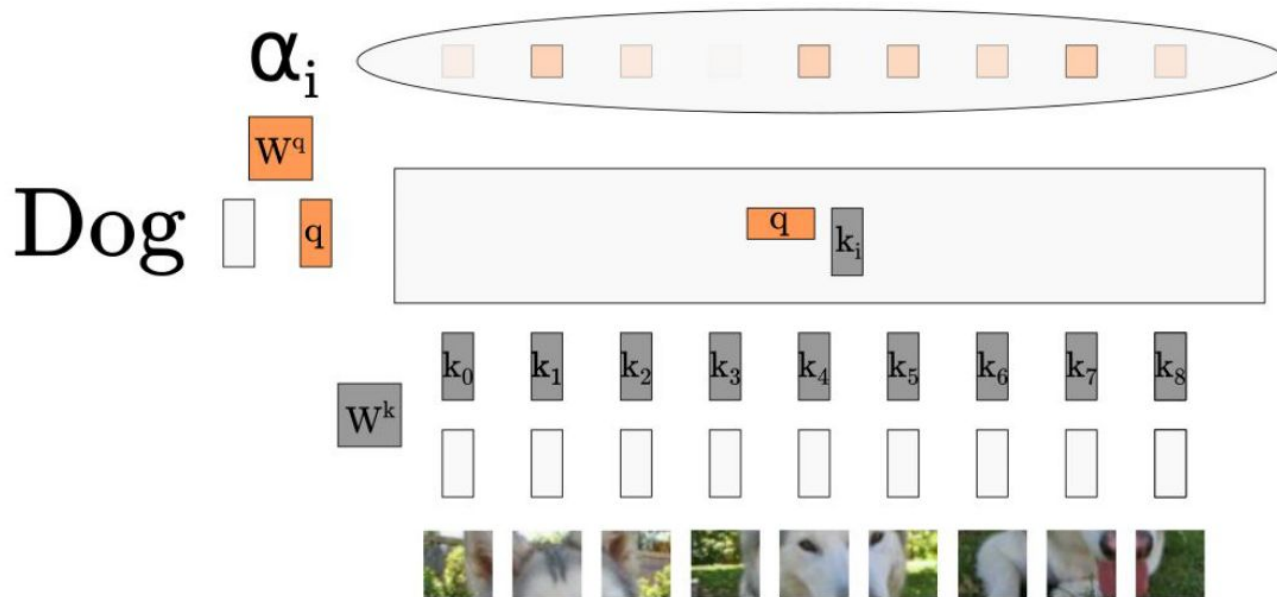


- Describe “Dog” by combining similar patch tokens

How?




- Describe “Dog” by combining similar patch tokens




- Find an alignment score through dot product


- Weighted sum of similar patches




- Weighted sum of similar patches




Scaled Dot-Product Attention




- Weighted sum of similar patches



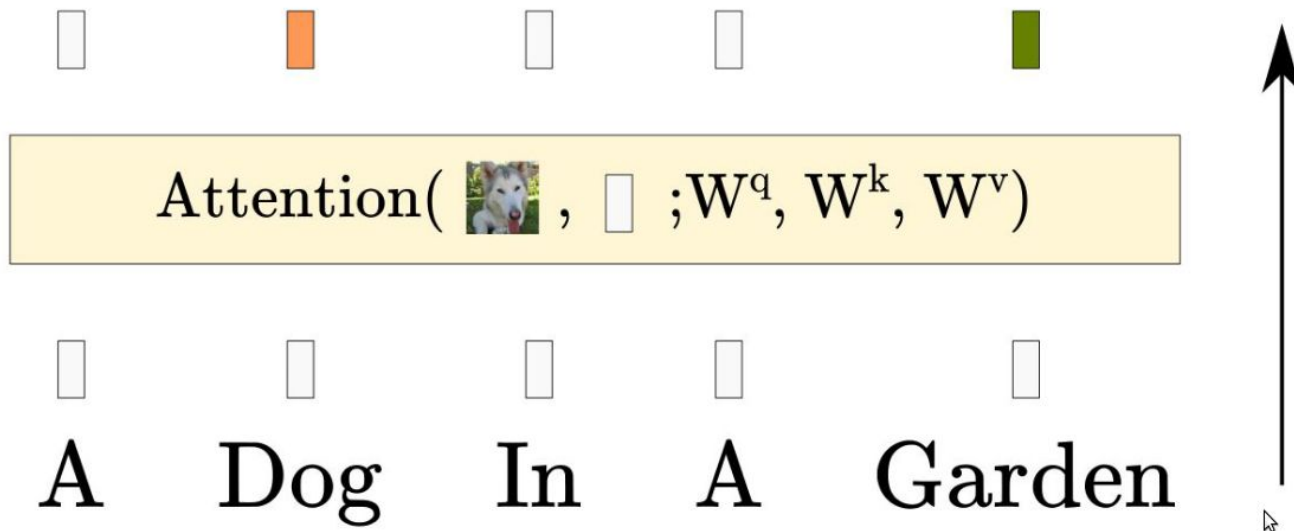
Scaled Dot-Product Attention



Multi-Head Attention



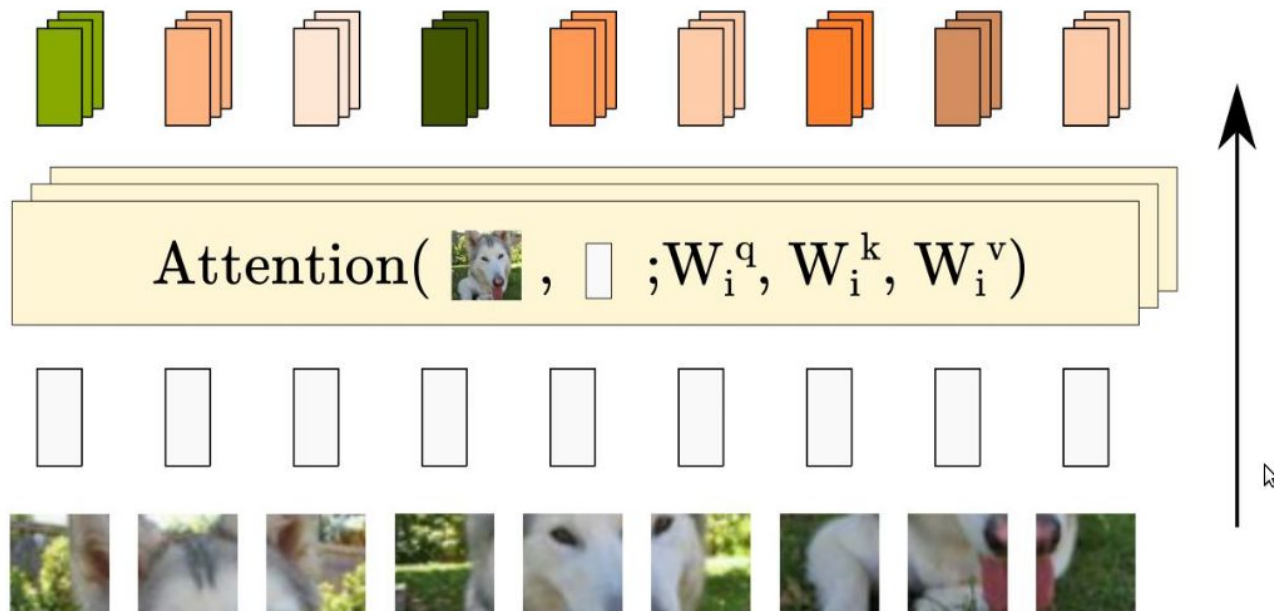
- Train up Weights to get good attention



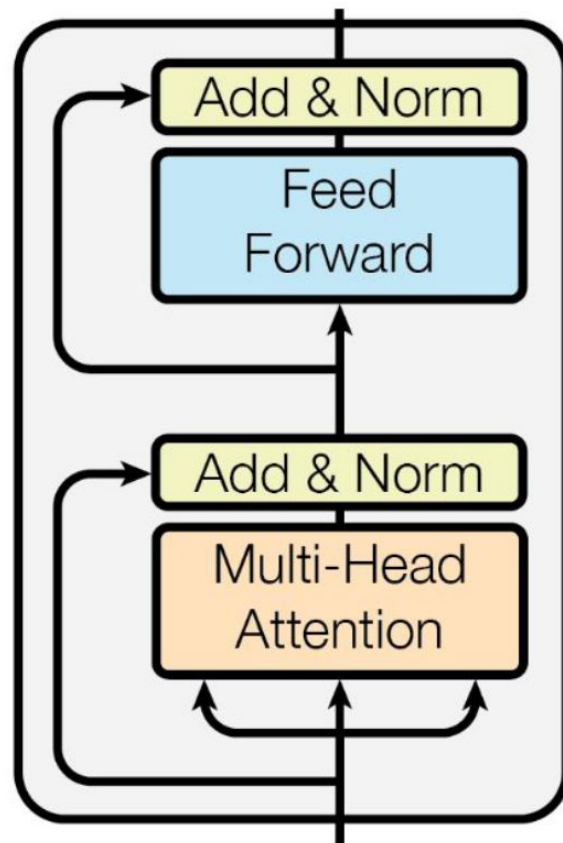
- You can compare a sequence to itself!



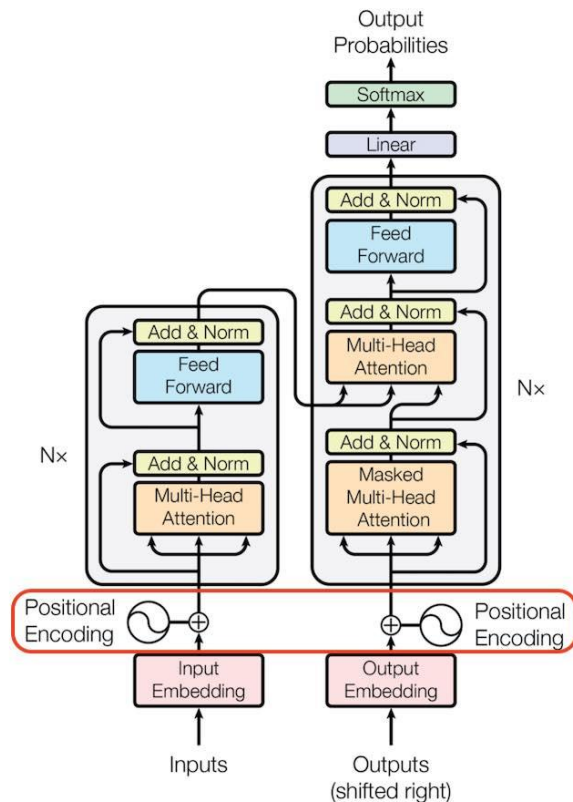
- And even have multiple interpretations



- Basic encoder block
 - MultiHead attention
 - Skip connection
 - Small MLP applied to each token
 - Skip connection
- Stacked in a transformer



- What is the order of the tokens?
 - Treated as a set
 - Permutation invariant
- How do keep positional info?
 - Masking
 - Add a positional encoding
 - Sine encoding
 - Learned encoding




- Attention: represent objects by similar things
- Mechanism
 - Compute alignment with product
 - Weighted sum of similar objects
 - Learn projection weights
- Transformer block
 - Attention layer
 - MLP on each token/object


What is great about attention?

- Attention can implement a lot of different operations
 - Little built-in bias
 - As opposed to CNNs
 - Adapts to data
 - Can change depending on the input
- A lot of work has been done to “discover” good operators
 - To little avail
 - But attention kind of does that!


- Possible to attend to all elements in the sequence
- (Very) loose approximation of dense layer



- Possible to use positional info to have convolutions
- Can be exactly approximated with ConViT (Touvron et al. 2021)

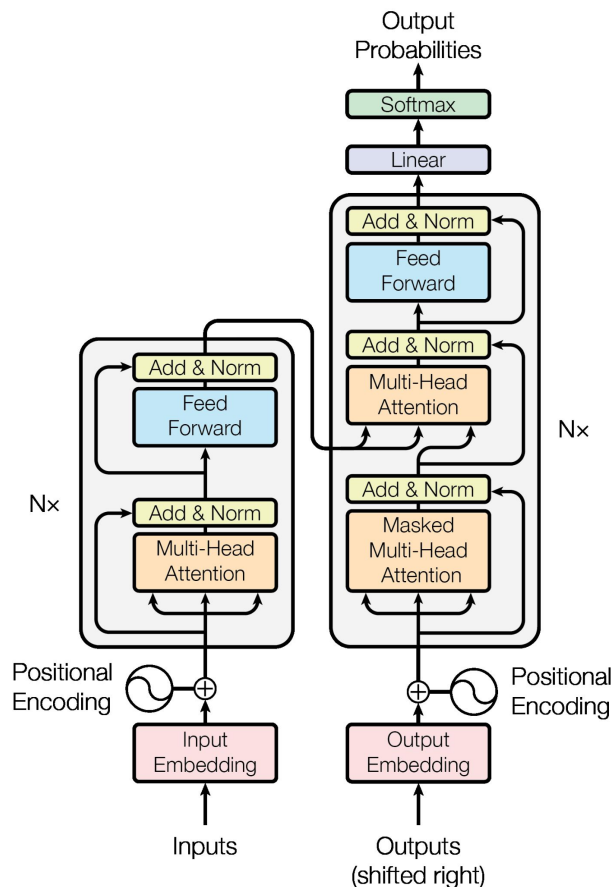


- Possible to just keep attention with past elements
- Default mode of transformer decoder



- Transformers are able to leverage large datasets
 - Because they can learn more adapted relations
 - Becoming more and more adopted
 - Scale very well
- Emerging as the dominant neural network type
- Drawbacks
 - Quadratic cost with the number of tokens
 - Need lots of data or strong regularization

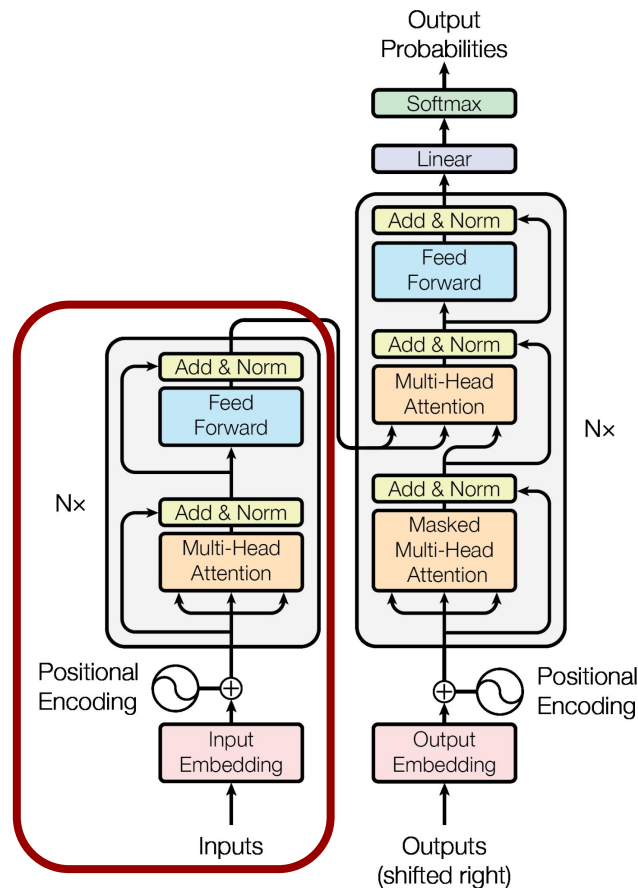
- Completely does away with recurrent units
 - Attention as a first class citizen!
 - Introduces element wise MLP for transform
- Transformer
 - *Transforms* the input throughout the layers
 - Also to blame for BERT, ELMO, DALL-E, ...



- Very strong results as soon as 2017!

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	

- Encoder yields strong features for the inputs
- Often used as standalone
 - No decoder
 - Linear layer or MLP directly plugged on encoder inputs
- Lots of applications!



- Detailed explanation of all transformers components
- Implement
 - Alignment score
 - Attention block
 - Encoder block
 - Transformer predictor
- Task: Reverse a sequence
 - With attention visualization!

