


Model-based statistical learning: Co-clustering with the latent bloc model

Vincent Vandewalle (vincent.vandewalle@univ-cotedazur.fr)

Msc 2 Data-Science & IA, 2024-2025

Co-clustering aims at performing simultaneous clustering of both rows and columns:



Source: [Christophe Biernacki, Julien Jacques, and Christine Keribin \(2022\)](#). “A Survey on Model-Based Co-Clustering: High Dimension and Estimation Challenges”. [In](#)

Bi-clustering, co-clustering and Latent Block Model (LBM)

- **Bi-clustering algorithms:** aim to detect homogeneous blocks within the data matrix which do not cover the entire matrix and which may overlap.
- **Co-clustering:** a specific bi-clustering model which assumes that all the individuals belong to one and only one row cluster, and symmetrically all the variables belong to only one column cluster.
- **Latent Block Model (LBM):** LBM is a model for performing a model-based co-clustering

See Sara C Madeira and Arlindo L Oliveira (2004). "Biclustering algorithms for biological data analysis: a survey". In: *IEEE/ACM transactions on computational biology and bioinformatics* 1.1, pp. 24–45 for more details on bi-clustering algorithms.

Questions on Model-Based Clustering (MBC)

- 1 Recall the principle of model-based clustering
- 2 For what type of data is it designed?
- 3 What is the link between the component of the mixture and the clusters?
- 4 How to select the number of clusters?
- 5 How can you compare two partitions when performing clustering?
- 6 Why using the rand index?
- 7 Why performing only clustering on rows, then on columns would not be sufficient to solve the co-clustering problem?

Questions on Model-Based Clustering (MBC)

- 1 Recall the principle of model-based clustering Model the distribution of the data as a mixture of distributions.
- 2 For what type of data is it designed? Any kind of data as soon as we are able to propose a model for the class specific density.
- 3 What is the link between the component of the mixture and the clusters? Each component is interpreted as a cluster
- 4 How to select the number of clusters? It can be selected by AIC, BIC or ICL
- 5 How can you compare two partitions when performing clustering? By using the Adjusted Rand Index ✓
- 6 Why using the rand index? It is invariant up to class permutation
- 7 ✓ Why performing only clustering on rows, then on columns would not be sufficient to solve the co-clustering problem? I allow to model the whole data matrix by a very sparse model.

The Latent Block Model (LBM) assumptions (1/2)

Data matrix \mathbf{x} ($n \times d$)

- \mathbf{x}_i : the row/individual number i
- \mathbf{x}^j : the column/variable number j of \mathbf{x}
- x_i^j : variable j of individual i

Partition of the rows \mathbf{z} ($n \times K$)

- $\mathbf{z} = (\mathbf{z}_1, \dots, \mathbf{z}_n)$
- $\mathbf{z}_i = (z_{i1}, \dots, z_{iK}) \in \{0, 1\}^K$
- $z_{ik} = 1$ if i belongs to row group k and 0 otherwise

Partition of the columns \mathbf{w} ($d \times L$)

- $\mathbf{w} = (\mathbf{w}_1, \dots, \mathbf{w}_d)$
- $\mathbf{w}_i = (w_{j1}, \dots, w_{jL}) \in \{0, 1\}^L$
- $w_{j\ell} = 1$ if variable \mathbf{x}^j belongs to column group ℓ and 0 otherwise

Main assumption: each point x_i^j is assumed to be independent given \mathbf{z}_i and \mathbf{w}_j (the knowledge of the block):

$$f(\mathbf{x}|\mathbf{z}, \mathbf{w}; \theta) = \prod_{k=1}^K \prod_{\ell=1}^L \prod_{i=1}^n \prod_{j=1}^d f(x_i^j; \alpha_{k\ell})^{z_{ik}w_{j\ell}}$$

with $f(\cdot; \alpha_{k\ell})$ the pdf associated to block $k\ell$ and parametrized by $\alpha_{k\ell}$.

The Latent Block Model (LBM) assumptions (2/2)

Moreover independence is assumed between all \mathbf{z}_i and \mathbf{w}_j :

$$f(\mathbf{z}, \mathbf{w}; \theta) = \prod_{i,k} \pi_k^{z_{ik}} \prod_{j,\ell} \rho_j^{w_{j\ell}}$$

with $\pi = (\pi_k)_k$ (the probabilities of each cluster in row), $\rho = (\rho_\ell)_\ell$ (the probabilities of each cluster in column). $\theta = (\pi, \rho, \alpha)$ groups all the parameters

Thus

$$f(\mathbf{x}, \mathbf{z}, \mathbf{w}; \theta) = \prod_{i,k} \pi_k^{z_{ik}} \prod_{j,\ell} \rho_j^{w_{j\ell}} \prod_{i,j,k,\ell} f(x_i^j; \alpha_{k\ell})^{z_{ik} w_{j\ell}}$$

Marginalizing over \mathbf{z} and \mathbf{w} (since they are not observed in practice ...), the pdf of \mathbf{x} is

$$f(\mathbf{x}; \theta) = \sum_{(\mathbf{z}, \mathbf{w}) \in \mathcal{Z} \times \mathcal{W}} \prod_{i,k} \pi_k^{z_{ik}} \prod_{j,\ell} \rho_j^{w_{j\ell}} \prod_{i,j,k,\ell} f(x_i^j; \alpha_{k\ell})^{z_{ik} w_{j\ell}}$$

with \mathcal{Z} (resp. \mathcal{W}) the set of all possible partitions of the rows (resp. the columns)

Choice of $f(\cdot; \alpha_{k\ell}$ according the type of data for x_i^j

- **Binary:** Bernoulli of parameter $\alpha_{k\ell}$
- **Categorical with r levels:** Multinomial distribution with parameters $\alpha_{k\ell} = (\alpha_{k\ell}^1, \dots, \alpha_{k\ell}^r)$
- **Count data:** Poisson distribution with parameter $\alpha_{k\ell}$
- **Continuous:** Normal distribution with parameters $\alpha_{k\ell} = (\mu_{k\ell}, \sigma_{k\ell}^2)$
- Can be extended to numerous other data types (ordinal, functional, textual, ...)

These models are very parsimonious even in high dimension!

ToDo : Count the number of parameters of the LBM for each data type

The observed log-likelihood is defined as:

$$\ell(\theta; \mathbf{x}) = \log f(\mathbf{x}; \theta) = \log \left(\sum_{(\mathbf{z}, \mathbf{w}) \in \mathcal{Z} \times \mathcal{W}} \prod_{i,k} \pi_k^{z_{ik}} \prod_{j,\ell} \rho_j^{w_{j\ell}} \prod_{i,j,k,\ell} f(x_i^j; \alpha_{k\ell})^{z_{ik} w_{j\ell}} \right)$$

- $\ell(\theta; \mathbf{x})$ requires the computation of $K^n L^d$ terms which correspond to all the possible configurations of unobserved labels \mathbf{z} and \mathbf{w} !
- The problem is a missing data problem thus possible to use the EM algorithm

$Q(\theta; \theta')$ the expectation of the completed log-likelihood

- $\ell_c(\theta; \mathbf{x}, \mathbf{z}, \mathbf{w})$ the completed likelihood
- $Q(\theta, \theta') = \mathbb{E}(\ell_c(\theta; \mathbf{x}, \mathbf{z}, \mathbf{w}); \mathbf{x}, \theta')$ the expectation of the completed log-likelihood given the current parameters θ'

EM algorithm starting from $\theta^{(0)}$ and loop until convergence

- Expectation (E) step: Computation of $Q(\theta; \theta')$
- Maximization (M) step: $\theta^{(q+1)} = \arg \max_{\theta} Q(\theta, \theta^{(q)})$

E step: computation of $Q(\theta, \theta^{(q)})$

The EM algorithm allows to increase the log-likelihood at each iteration: $\ell(\theta^{(q+1)}) \geq \ell(\theta^{(q)})$ and thus to converge to a local maximum of the likelihood

$$\ell_c(\theta; \mathbf{x}, \mathbf{z}, \mathbf{w}) = \sum_k \left(\sum_i z_{ik} \right) \log \pi_k + \sum_\ell \left(\sum_j w_{j\ell} \right) \log \rho_\ell + \sum_{i,j,k,\ell} \log f(x_i^j; \alpha_{k\ell})$$

Thus by taking the conditional expectation, we get:

$$\begin{aligned} Q(\theta, \theta^{(q)}) &= \sum_{i,k} p(z_{ik} = 1 | \mathbf{x}, \theta^{(q)}) \log \pi_k + \sum_{j,\ell} p(w_{j\ell} = 1 | \mathbf{x}, \theta^{(q)}) \log \rho_\ell \\ &\quad + \sum_{i,j,k,\ell} p(z_{ik} w_{j\ell} = 1 | \mathbf{x}; \theta^{(q)}) \log f(x_i^j; \alpha_{k\ell}) \end{aligned}$$

Let $s_{ik}^{(q)} = p(z_{ik} = 1 | \mathbf{x}; \theta^{(q)})$, $t_{j\ell}^{(q)} = p(w_{j\ell} = 1 | \mathbf{x}; \theta^{(q)})$ and $p(z_{ik} w_{j\ell} = 1 | \mathbf{x}; \theta^{(q)})$. All these computations are intractable due to dependence structure in the model.

Question: Assume that you would know these intractable quantities, how would perform the M-step?

- Variational approach: Constrain the joint probability to satisfy the relation

$$p(\mathbf{z}, \mathbf{w}|\mathbf{x}; \theta) \approx p_z(\mathbf{z}|\mathbf{x}; \theta)p_w(\mathbf{w}|\mathbf{x}; \theta)$$

where p_z and p_w are chosen to provide the closest approximation of $p(\mathbf{z}, \mathbf{w}|\mathbf{x}; \theta)$ while still being computable. The algorithm maximizes an evidence lower bound (ELBO)

$$\ell(\theta; \mathbf{x}) \geq \mathcal{F}(\theta; \mathbf{x}) = \max_{p_z, p_w} (\ell_c(\theta; \mathbf{x}, \mathbf{z}, \mathbf{w}) - \log(p_z(\mathbf{z})p_w(\mathbf{w})))$$

this algorithm is called VEM as variational EM

- SEM algorithm : alternates the following steps: simulate $\mathbf{z}|\mathbf{x}, \mathbf{w}; \theta$ and then $\mathbf{w}|\mathbf{x}, \mathbf{z}; \theta$. Then update θ given the simulated classes \mathbf{z} and \mathbf{w}

Estimating and evaluation of the rows and the columns clusters

Estimation

- VEM : based on $p_z(\mathbf{z}|\mathbf{x}; \hat{\theta})$ and $p_w(\mathbf{w}|\mathbf{x}; \hat{\theta})$ at the last iteration
- SEM: Based on sampling $(\mathbf{z}, \mathbf{w})|\mathbf{x}; \hat{\theta}$ by a Gibbs sampler, then estimate $(\hat{\mathbf{z}}, \hat{\mathbf{w}})$ by the mode of the marginal sampled distribution.

Evaluation

- ARI: Adjusted Rand Index / For the rows and columns respectively
- CARI: Co-clustering ARI developed for co-clustering